

Enterprise inter-VLAN traffic analysis

Abstract—In this short paper, we analyze the traffic captured in the enterprise network. For the convenience, we have chosen VLAN router as our vantage point due to the simplicity of the traffic capturing process. Such a choice of a vantage point prevented us from installing multiple sniffers in the network, but since our goal was to investigate the interactions of the computers in different VLANs such setup was natural.

With this respect in the proceeding paragraphs we show the results of our quantitative study and discuss key findings.

I. INTRODUCTION

Analyzing the traffic in the company on regular basis can be beneficial for several reasons: (i) we can find strange behaviour of computers in the network, (ii) we can find the bottlenecks, (iii) we can detect misconfigurations in networked devices. In this short paper we present the analysis traffic which was captured in a small enterprise.

II. DATA COLLECTION METHODOLOGY

We begin this section with a description of the enterprise network, with the details of data capturing process to follow. According to Cisco documentation [1], our network was two-tier medium-sized enterprise network comprising slightly more than 200 workstations, one distribution switch (Cisco Catalyst G3560 series layer 3 switch), one router (the router was running Linux operating system) playing the role of a *VLAN router*, 6 access switches from various vendors (1 DLink DES3200, 1 Cisco 2690, 2 DLink DES1226 and 2 XNET SH9024 switches) and also multiple, non-managed switches and one wireless access point all being connected to access switches.

Since we were unable to analyze fully the wiring of the access (especially those that were spread around the building and connected directly to access switches) and distribution switches, we could basically only guess (based on the limited information we had at our disposal) that there were no redundant links between access and distribution switches, also there were no redundant links between distribution switch and VLAN router. Essentially, the switches were merely connected in a tree like topology: each access switch was connected to one distribution switch, and the distribution switch was connected to a single router over a gigabit Ethernet trunk link.

The entire network was partitioned into 25 broadcast domains (VLANs), with native VLAN (VLAN that does not have tagging) being reserved for infrastructure management, although not all the switches were using native VLAN for management purposes. For brevity and privacy reasons we omit the discussion of network addressing in full detail, but rather mention that all VLANs were assigned addresses from class C private network from 192.168.0.0 to 192.168.24.0, whereas the last usable address in each subnetwork was used

as the address for the default gateway. All hosts in the network used static address assignment schema.

To capture the traffic it was natural to place packet capturing at the central router, rather than installing it on multiple switches. With this respect, we configured sniffer at the router to capture the frames on a gigabit trunk interface. We have used *tcpdump* to capture all traffic on the trunk link of a router - the link which was connected to the distribution switch. The data collection period lasted for slightly more than 2.5 hours, from 06 : 23 UTC to 09 : 03 UTC and $19.1 \cdot 10^6$ frames were collected in total. Once the data was collected, a single *pcap* file was created. We then started to preprocess the captured data.

III. DATA PROCESSING AND BASIC RESULTS

To perform the preprocessing we have extracted the binary data and stored it in *pdml* (Packet Description Markup Language) format - a special XML format which is used to describe the contents of the binary representation of frames and packets.

Our next step was to extract needed fields from the XML file and store the extracted data in the format which was simpler for further analysis. Thus, we have chosen JSON (JavaScript Object Notation) format. From PDML we have extracted the following fields:

- Ethernet source and destination MAC addresses and type field
- VLAN identifier and type field
- From IP header we have extracted source and destination addresses, version number, length, protocol type, and TTL.
- Basic information from GRE and PPP headers (since clients accessing the Internet were using these protocols)
- From TCP header we have selected source and destination ports, stream identifier, length, sequence and acknowledgment numbers, as well as flags and window size.
- From UDP header source and destination ports, as well as datagram length information were extracted.

Our next step in preprocessing the data was to extract non IPv4 frames - frames for which neither Ethernet type nor VLAN type field were equal to **0x00008000** (this dataset includes packets from all VLANs as well as packets from native VLAN). Once the data was extracted we binned the frames according to protocol types. In Table I we present the summary results for this data.

The next step in preprocessing of the data was to exclude the frames for the native VLAN from the traces. Thus, we have

TABLE I: Distribution of non IPv4 frames

Protocol	Number of frames	Fraction (%)
ARP	136857	48.8
Cisco Shared Spanning Tree Protocol	119385	42.5
IPv6	17034	6.1
Spanning Tree Protocol (IEEE 802.1D)	4774	1.7
IPX	1164	0.4
Cisco Loop	957	0.3
Cisco CDP/VTP	510	0.2

TABLE II: Distribution of protocols in native VLAN

Protocol	Number of frames	Fraction (%)
Cisco Shared Spanning Tree Protocol	4774	40.5
Spanning Tree Protocol (IEEE 802.1D)	4774	40.5
Cisco Loop protocol	957	8.1
IP (ICMP only)	640	5.4
Cisco CDP/VTP	510	4.3
ARP	128	1.1

TABLE III: Frames with invalid MAC addresses

Source MAC	Destination MAC	Source IP	Destination IP	Manufacturer
30:f9:ed:41:a6:01	00:c0:ee:9a:5a:85	192.168.5.10	192.168.5.151	Sony/KYOCERA
00:30:05:c2:b7:ff	00:17:e8:03:a0:7b	192.168.18.7	192.168.18.34	Fujitsu/KYOCERA
00:15:58:67:5f:14	00:c0:ee:9a:5a:85	192.168.5.3	192.168.5.151	FOXCONN/KYOCERA

filtered out the frames for which Ethernet type was not equal to **0x00008100**. It turned out that the fraction of frames without VLAN tag was rather small and constituted only 0.06% of the total number of frames in the trace. Since native VLAN was used only for management purposes we excluded it from further analysis. However, in Table II we show the distribution of packet types seen in the native VLAN.

Instead, next we turned our attention to frames with invalid MAC addresses. By invalid MAC address we mean those MAC addresses which are not multicast or broadcast addresses received at the trunk interface of the router with destination unicast MAC addresses not of router's own MAC addresses. Upon filtering the data we have found that there were two such destination unicast MAC addresses. We hypothesize that such frames could have been received at the trunking interface due to the following reasons. The MAC address table of the switches (i) did not contain a record for the destination MAC address and so it was flooded to all ports except the port from which the frame was received; (ii) wrong mapping for the destination MAC address and outgoing interfaces could have existed and so it was forwarded into the wrong port of the switch and thus was received at the trunk port of the router. The two invalid MAC addresses we have discovered are: **00:17:e8:03:a0:7b** and **00:c0:ee:9a:5a:85**. We leave this investigation to the network administrators.

Our next step of data preprocessing was to find for every packet being forwarded a corresponding pair: Note, our trace contained two copies of a packet for which the source and destination IP addresses where within the **192.168.0.0/16** subnetwork with a difference that TTL was reduced by one for one of the packets, and the Ethernet header was recalculated. Thus, by filtering out the duplicates we were able to exclude the possibility of overcounting the number of bytes carried in TCP and UDP streams, and other transport protocols. To perform this filtering step, we merely found all packets whose source and destination addresses where within the subnetwork

192.168.0.0/16 and filtered out frames for which source MAC address was not equal to the MAC address of the VLAN server (note, all VLAN interfaces were assigned the same MAC address). Thus, effectively leaving only one copy of the packet being forwarded between the subnetworks in the trace. Our resulting trace was reduced by 35.4%, and now contained 12190001 frames.

Once again we have searched for frames with invalid MAC addresses (we have introduced the term invalid MAC addresses in the previous paragraphs). In Table III we list these MAC addresses as well as corresponding IP addresses found in such frames, and manufacturer's name. Obviously, such discrepancy is strange.

For the cleaned data, our first step was to take a look at the distribution of TCP and UDP applications and the number of bytes these applications transmitted. Thus, in Table IV we show the distribution of 9 most used UDP and TCP applications. We examined both inter-VLAN and Internet traffic. We observed that while most of hosts were not actively connected to the Internet, using the VPN connections, 15366 connection attempts were made towards the Internet (only TCP SYN packet was captured, which was dropped by the firewall). This could be an indication that quite a lot of connections are made in fact in a stealth mode on users' computers. Next, we have computed the adjacency graph for computers which were interacting over TCP in the local network (an interaction here means that at least one TCP/UDP connection between the pair of computers existed). As it was expected, only few computers (servers) were accepting all TCP connection, with a small fraction of computers were interacting between each other. Upon expecting closely the logs we have found few interesting things: (i) at least one machine had misconfigured address (it was using self generated IP addresses with the prefix 169.254.171.0/24); (ii) only two machines were interacting with non-server computers 192.168.23.13 and 192.168.3.6 using ports 139 and 445.

TABLE IV: Distribution of top applications used in the network

TCP port	Frequency	Application	UDP port	Frequency	Application
13000	11599	Kaspersky	53	27960	DNS
443	10128	HTTPS	137	6589	NetBIOS
88	9700	Kerberos	389	1350	LDAP
80	6938	HTTP	15000	804	Kaspersky Network Agent
445	5581	Microsoft SMB	88	166	Kerberos
135	940	MS RPC	123	146	NTP
389	880	LDAP	13000	50	Kaspersky
49155	694	Microsoft-DC	138	21	NetBIOS
139	682	Netbios	443	16	QUIC

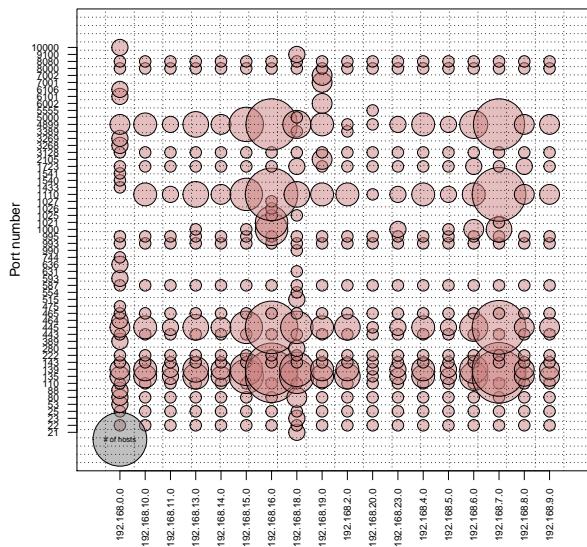


Fig. 2: Distribution of ports used in the network

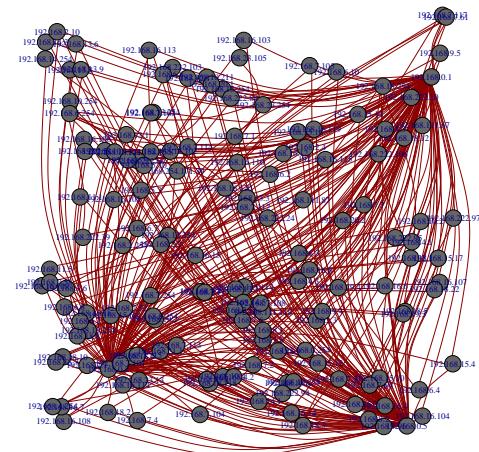


Fig. 1: Interaction between computers in different VLANs over TCP

Next we computed the distribution of the packet sizes for packets with valid IP addresses. Interestingly there were frames which were larger than maximum allowed Ethernet frames size (so called jumbo frames). Around 10% of the IPv4 packet were larger than 1500 bytes. Finally, we have built the distribution of destination ports used in the packets. In Figure 2 we show this distribution.

IV. CONCLUSIONS

In this short paper we have played a bit with the packets which were captured in a small enterprise. Our primary goal was to analyse the interactions of computers and build the statistics for the traffic which was captured for several hours. Our key findings are the following: (i) major traffic in the network is HTTPS and HTTP, (ii) antivirus solutions consume considerable amount of bandwidth, (iii) computers in the network mainly interact with the default gateway and few servers, such as NFS server and mail server.

REFERENCES

- [1] W. Odom. *CCENT/CCNA ICND1 100-105 Official Cert Guide*. 2016.